

Pre-flight checklist — run this 30 min before Tom joins

Brazil time: do this at 09:30 for a 10:00 call.

30 minutes out

- **Laptop plugged in.** GPU jobs will chew battery.
- **Close Slack, Discord, email, notifications.** Tom sees your screen.
- **Open Tom's Lab** via `Run Tom's Lab.bat`. **Leave the disclaimer dialog up** — Tom sees it cold during the demo (see 04_demo_script.md at 00:00). Do NOT click through pre-call.
- **Reset the disclaimer flag** if you've previously accepted it on this machine, so it pops up live for Tom: ask the dev to clear `disclaimer_accepted` in the settings table.
- **Check AI providers are alive:**
 - Settings → Gemini API key set? Click *Test connection* → expect green tick.
 - Ollama running? (taskbar icon) — it's the offline fallback.
 - If either fails, fix it now. A "ProviderUnavailable" error in front of Tom kills the demo.

20 minutes out

- **Warm Semantic search.** Switch mode to *Semantic*, type `VPOC`, hit enter. First run will trigger an embedding build if one isn't cached — wait for the status bar to finish. Subsequent searches will be instant.
- **Warm Visual search.** Switch mode to *Visual*, type `chart with three peaks`. Confirms CLIP is loaded and Gallery will have real hits.
- **Populate Gallery.** Run a broad keyword search (e.g. `Tom`) so the Gallery tab isn't empty when you click it. Empty Gallery looks like a broken app.
- **Pre-open Ask Tom** and ask one warm-up question so the first-call latency is paid before Tom sees it. Suggested: *"Explain Naked VPOC (NVPOC) and why Tom watches for it."* — this is also your cold-open, so the second run will feel snappy.

10 minutes out

- **Window size.** Maximise Tom's Lab. Don't share a tiny window.
- **Font size.** If you're sharing to a call, the default reads small on a compressed video feed. Either bump your OS display scale to 125% or verbally say you'll zoom in when needed.

- **Docs A & B sent.** Email Tom both PDFs before the call starts, subject line: *"Tom's Lab — quick read before our call"*. Don't rely on him reading them live.
- **Demo script open on your phone or a second monitor**, not the screen you're sharing.
- **Glass of water.** It's a 45-min talk, you'll dry out.

2 minutes out

- **Breathe.** Tom didn't answer your questions before — that's not a "no", that's a busy subject-matter expert. Today you're making it easy for him.
 - Confirm your tab is **Feed**, scroll position **top**, search bar **empty**. Clean start.
-

If something breaks mid-demo

Symptom	What to say	What to do
Ask Tom hangs	<i>"The model is slow today; I'll skip ahead and come back."</i>	Move on; don't cancel mid-answer — Gemini usually finishes in 5–15 s.
Provider error	<i>"AI providers are optional; the keyword and visual search still work entirely offline — let me show those."</i>	Switch to Keyword/Visual. The non-AI features are the safe ground.
Empty Gallery	<i>"Haven't run a search yet this session."</i>	Run a broad keyword search, then return to Gallery.
Transcription pill shows up	<i>"That's TomTube indexing your YouTube in the background — no need to watch it."</i>	Ignore it.
App crashes	<i>"Let me restart — happens once in a while, data's safe on disk."</i>	Relaunch via the batch file; all state (bookmarks, favourites, index) is persisted.

After the call

- Write Tom's answers into 02_questions_for_tom.md while they're fresh.
- Note any new questions he raised — those become your next list.
- Save any off-script reactions ("I'd want X to not be there") — those become feedback-memory entries.